

The Conjunctive Disjunctive Graph Node Kernel

Dinh Tran Van, Alessandro Sperduti

Department of Mathematics, Padova University, Trieste, 63, 35121 Padova, Italy

Fabrizio Costa

Department of Computer Science, University of Exeter Exeter EX4 4QF, UK

Abstract

Gene-disease associations are inferred on the basis of similarities between the proteins encoded by genes. Biological relationships used to define similarities range from interacting proteins, proteins that participate in pathways and protein expression profiles. Though graph kernel methods have become a prominent approach for association prediction, most solutions are based on a notion of information diffusion that does not capture the specificity of different network parts. Here we propose a graph kernel method that explicitly models the configuration of each gene’s context. An empirical evaluation on several biological databases shows that our proposal achieves state-of-the-art results.

Keywords: Graph node kernels, graph decomposition, disease gene prioritization

1. Introduction

Predictive systems for gene-disease associations are often based on a notion of similarity between genes. A common strategy is to encode relations between genes as a network and use graph based techniques to make useful inferences.

5 In the last decades, a number of graph kernel methods have been proposed that directly exploit transitive properties in biological networks. The prototypical method is the Diffusion kernel (DK) [1] inspired by the heat diffusion phenomenon. The key idea is to allow a given amount of *heat* placed on nodes

to *diffuse* through the edges. The similarity between two nodes v_i, v_j is then
10 defined as the amount of heat starting from v_i and reaching v_j within a given
time interval. In DK the heat flow is proportional to the number of paths
connecting two nodes, which introduces an undesired bias that penalize periph-
eral nodes w.r.t. central ones. This problem is tackled by a modified version
of DK called Markov exponential diffusion kernel (MED) [2] where a Markov
15 matrix replaces the Laplacian matrix. Another kernel called Markov diffusion
kernel (MD) [3], exploits instead the notion of *diffusion distances*, a measure of
similarity between patterns of heat diffusion. The Regularized Laplacian ker-
nel (RL) [4] represents instead a normalized version of the random walk with
restart model and defines the node similarity as the number of paths connecting
20 two nodes with different lengths. All these approaches can be applied to dense
networks with high degree nodes. A drawback of these approaches is however
their relatively low discriminative capacity. This is in part due to the fact that
information is processed in an additive and independent fashion which prevents
them from accurately modeling the configuration of each gene’s context. To
25 address this issue here we propose to employ a *decompositional* graph kernel
(DGK) [5] technique in which the similarity function between graphs can be
formed by decomposing each graph into subgraphs and by devising a valid local
kernel between the subgraphs. To exploit its higher discriminative capacity we
first decompose the network in a collection of connected sparse graphs and then
30 we develop a suitable kernel, that we call Conjunctive Disjunctive Node Kernel
(CDNK).

2. Material and methods

We start from the type of similarity notion computed by decomposition
kernels between graph instances and adapt it to express the similarity between
35 nodes in a single network. In this work we use three key ideas: 1) genes are
labeled using their functional profile, 2) we transform the network to distinguish
highly connected components from sparsely connected ones, and 3) we transform

the neighborhood of each gene in a sparse high dimensional vector that can be easily processed by standard machine learning techniques such as SVMs.

40 2.1. Gene Labeling

Gene-disease associations networks typically represent genes as nodes labeled with a gene identifier. Here we take a different approach and use the node labels to encode abstract information about the genes. In this way downstream machine learning algorithms can generalize from similar examples and allow the
45 identification of overlooked but related genes. We experiment with two types of information: 1) topological information and 2) functional information.

Topological information is simply based on the connectivity degree of the gene. The idea is that genes that have the same number of connections are more similar than genes with a different connectivity.

50 Functional information is based on the *Gene Ontology* [6] resource. We use the ontology to construct binary vectors representing a bag-of-words encoding for each gene (i.e. if a GO-term is associated with the gene). The resulting representations are then clustered so that genes with similar description profiles receive the same class identifier as label. All subsequent graph processing
55 (including graphs isomorphism detection) takes this discretized labeling into account.

2.2. Vec labels

2.3. Network Decomposition

In gene-disease associations networks it is not uncommon to find nodes with
60 high degrees. Unfortunately these cases cannot be effectively processed by a neighborhood based decomposition kernel since these are based on the notion of exact matches and the probability of finding identical neighborhoods decreases exponentially as the degree increases. This means that in a finite network it quickly becomes impossible to find any match and hence learn or generalize at
65 all. As an alternative, we propose a procedure to "sparsify" the network that is observed by the neighborhood kernel. In reality we do not alter the cardinality of

the edge set, but rather mark the edges with special attributes that will inform the kernel computation. The result is a procedure that decomposes the network in a linked collection of sparse sub-networks where each node has a reduced
70 connectivity when considering the edges of a specific type. However the other edges are still available to connect the various sub-networks. We distinguish two types of edges: *conjunctive* and *disjunctive* edges. Nodes linked by conjunctive edges are going to be used jointly to define the notion of context and will be visible to the neighborhood graph kernel. Nodes linked by disjunctive edges are
75 instead used to define features based only on the pairwise co-occurrence of the genes at the endpoints and are processed by our novel kernel.

Definitions. We represent a problem instance as a graph $G = (V, E)$ where V is the set of nodes and E is the set of links. We define the *distance* $\mathcal{D}(u, v)$ between two nodes u and v , as the number of edges on the shortest path between
80 them. The *neighborhood* of a node u with radius r , $N_r(u) = \{v \mid d(u, v) \leq r\}$, is the set of nodes at distance no greater than r from u . The corresponding *neighborhood subgraph* \mathcal{N}_r^u is the subgraph induced by the neighborhood (i.e. considering all the edges with endpoints in $N_r(u)$). The *degree* of a node u , $d(u) = |\mathcal{N}_1^u|$, is the cardinality of its neighborhood.

85 *Iterative k-core decomposition* [7]: The node set is partitioned in two groups on the basis of the degree of each node w.r.t. a threshold degree D , the first part contains all nodes with degree smaller or equal than D and the second part the others. The node partition is used to induce the "conjunctive" vs "disjunctive" notion for the edge partition: edges that have both endpoints in the same part
90 (be it the first or the second) are marked as conjunctive, otherwise they are marked as disjunctive. We apply the k-core decomposition iteratively. At each iteration we consider only the graph induced by the conjunctive edges. We stop iterating the decomposition when no node has a degree ¹ greater than D . Note that this decomposition does not alter the cardinality of the edge set, it is simply
95 a procedure to mark each edge with the attribute conjunctive or disjunctive.

¹The degree is defined by only considering incident conjunctive edges.

Clique decomposition [8]: To model the notion that nodes in a clique are tightly related, we summarize the whole clique with a new 'representative' node. All the cliques (completely connected subgraphs) with a number of nodes greater than a threshold size C are identified. The endpoints of all edges incident on the clique's nodes are transferred to the representative node. Disjunctive edges are introduced to connect each node in the clique to the representative. Finally all edges with both endpoints in the clique are removed.

In our work a network is transformed by applying first the iterative k-core decomposition and then the clique decomposition.

2.4. Node Graph Kernels

We start from the Neighborhood Subgraph Pairwise Distance Kernel (NSPDK) [9] and adapt it to express the similarity between nodes in a single network. The key idea in NSPDK is to decompose graphs in small fragments and count how many pairs of fragments are shared between two instances. We introduce two improvements: 1) we partition the features according to the individual node's neighborhood, and 2) we introduce a distinction between "disjunctive" and "conjunctive" edges.

2.4.1. The Neighborhood Subgraph Pairwise Distance Kernel

The NSPDK is an instance of convolution kernel [5] where given a graph $G \in \mathcal{G}$ and two rooted graphs A_u, B_v , the relation $R_{r,d}(A_u, B_v, G)$ is true iff $A_u \cong \mathcal{N}_r^u$ is (up to isomorphism \cong) a neighborhood subgraph of radius r of G and so is $B_v \cong \mathcal{N}_r^v$, with roots at distance $\mathcal{D}(u, v) = d$. We denote R^{-1} as the inverse relation that returns all pairs of neighborhoods of radius r at distance d in G , $R_{r,d}^{-1}(G) = \{A_u, B_v | R_{r,d}(A_u, B_v, G) = \text{true}\}$. The kernel $\kappa_{r,d}$ over $\mathcal{G} \times \mathcal{G}$, counts the number of such fragments in common in two input graphs:

$$\kappa_{r,d}(G, G') = \sum_{\substack{A_u, B_v \in R_{r,d}^{-1}(G) \\ A'_{u'}, B'_{v'} \in R_{r,d}^{-1}(G')}} \mathbf{1}_{A_u \cong A'_{u'}} \cdot \mathbf{1}_{B_v \cong B'_{v'}},$$

where $\mathbf{1}_{A \cong B}$ is the *exact matching function* that returns 1 if A is isomorphic to B and 0 otherwise. Finally, the NSPDK is defined as $K(G, G') = \sum_r \sum_d \kappa_{r,d}(G, G')$,

where for efficiency reasons, the values of r and d are upper bounded to a given
125 maximal r^* and d^* , respectively.

2.4.2. The Conjunctive Disjunctive Node Kernel

We extend NSPDK and define a node kernel $K(G_u, G_{u'})$ between two copies
of the same network G where we distinguish the nodes u and u' respectively.
The idea is to define the features of a node u as the subset of NSPDK features
130 that always have the node u as one of the roots. In addition we distinguish
between two types of edges, called *conjunctive* and *disjunctive* edges. When
computing distances to induce neighborhood subgraphs, only conjunctive edges
are considered. When choosing the pair of neighborhoods to form a single
feature, we additionally consider roots u and v that are not at distance d but
135 such that u is connected to w via a disjunctive edge and such that w is at
distance d from v . In this way disjunctive edges can still allow an *information*
flow even if their endpoints are only considered in a pairwise fashion and not
jointly.

Formally, we define two relations: the *conjunctive relation* $R_{r,d}^\wedge(A_u, B_v, G_u)$
140 identical to the NSPDK relation $R_{r,d}(A_u, B_v, G)$, and (ii) $\mathcal{D}(u, v) = d$; the
disjunctive relation $R_{r,d}^\vee(A_u, B_v, G_u)$ is true iff (i) $A_u \cong \mathcal{N}_r^u$ and $B_v \cong \mathcal{N}_r^u$ are
true, (ii) $\exists w$ s.t. $\mathcal{D}(w, v) = d$, and (iii) (u, w) is a disjunctive edge. We define
 $\kappa_{r,d}$ on the inverse relations $R_{r,d}^{\wedge^{-1}}$ and $R_{r,d}^{\vee^{-1}}$

$$\kappa_{r,d}(G_u, G_{u'}) = \sum_{\substack{A_u, B_v \in R_{r,d}^{\wedge^{-1}}(G_u) \\ A_{u'}, B_{v'} \in R_{r,d}^{\wedge^{-1}}(G_{u'})}} \mathbf{1}_{A_u \cong A_{u'}} \cdot \mathbf{1}_{B_v \cong B_{v'}} + \sum_{\substack{A_u, B_v \in R_{r,d}^{\vee^{-1}}(G_u) \\ A_{u'}, B_{v'} \in R_{r,d}^{\vee^{-1}}(G_{u'})}} \mathbf{1}_{A_u \cong A_{u'}} \cdot \mathbf{1}_{B_v \cong B_{v'}}.$$

145 The CDNK is finally defined as $K(G_u, G_v) = \sum_r \sum_d \kappa_{r,d}(G_u, G_v)$, where once
again for efficiency reasons, the values of r and d are upper bounded to a given
maximal r^* and d^* .

3. Experimental

We perform an empirical evaluation of the predictive performance of several
150 kernel based methods on two of the databases used in [2].

BioGPS: A gene co-expression network is constructed from BioGPS dataset, which contains 79 tissues, measured with the Affymetrix U133A array. Edges are inserted when the pairwise Pearson correlation coefficient (PCC) between genes is larger than 0.5.

155 **Pathways:** Pathway information is retrieved from KEGG, Reactome, PharmGKB and the Pathway Interaction Database. If a couple of proteins co-participate in any pathway, the two corresponding genes are linked.

To evaluate the performance of graph node kernels we analyze the *gene prioritization*, i.e. given a set of genes known to be associated to a given disease, 160 gene prioritization is the task to rank the candidate genes based on their probabilities to be related to that disease. Similar to the evaluation process used in [2], we choose 12 diseases [10] with at least 30 confirmed genes. For each disease, we construct a positive set \mathcal{P} with all confirmed disease genes, and a negative set \mathcal{N} which contains random genes associated at least to one disease 165 class which is not related to the class that is defining the positive set. In [2] the ratio between the dataset sizes is chosen as $|\mathcal{N}| = \frac{1}{2}|\mathcal{P}|$. The predictive performance of each method is evaluated via a leave-one-out cross validation: one gene is kept out in turn and the rest are used to train an SVM model. We compute a decision score q_i for the test gene g_i as the top percentage value of 170 score s_i among all candidate gene scores. We collect all decision scores for every gene in the training set to form a global decision score list on which we compute the AUC ROC.

Model Selection: The hyper parameters of the various methods are set using a k-fold on a dataset set that is then never used in the predictive performance estimate. We try the values for diffusion parameter in DK and MED in 175 $\{10^{-3}, 10^{-3}, 10^{-2}, 10^{-1}\}$, time steps in MD in $\{1, 10, 100\}$ and RL parameter in $\{1, 4, 7\}$. For CDNK, we try for the degree threshold values in $\{10, 15, 20\}$, clique size threshold in $\{4, 5\}$, maximum radius in $\{1, 2\}$, maximum distance in $\{2, 3, 4\}$. Finally, the regularization trade off parameter C for the SVM is 180 searched in $\{10^{-5}, 10^{-4}, 10^{-3}, 10^{-2}, 10^{-1}, 1, 10, 10^2\}$.

4. Results and discussion

Table 1 and 2 show the AUC performance of the models trained by using different graph node kernels on 11 gene-disease association problems using the BioGPS and Pathways datasets to materialize the gene relation network. We
185 note that CDNK ranks first in 7 out of 11 cases using both networks. CDNK is the top performant kernel also when considering the average AUC ROC and the average rank with 73.3/2.0, 76.5/1.8, with a difference w.r.t. the second best of 5.5% and 1% on BioGPS and Pathways, respectively. MED and RL show similar and moderate results with small gap between them. DK and MD
190 exhibit modest performance on average and are ranked last in many cases: 7 times out of 11 for MD in BioGPS and 10 out of 11 for DK in Pathways. Low performance cases include networks that have high *average* node degree which is likely to yield very sparse and fragmented network decompositions.

5. Conclusions

195 We have shown how decomposing a network in a set of connected sparse graphs allows us to take advantage of the discriminative power of CDNK, a novel decomposition kernel, to achieve state-of-the-art results. In future work we will investigate how to 1) decompose networks in a data driven way and 2) extend the CDNK approach to gene-disease association problems exploiting
200 multiple heterogeneous information sources.

References

- [1] R. I. Kondor, J. Lafferty, Diffusion kernels on graphs and other discrete input spaces, in: ICML, Vol. 2, 2002, pp. 315–322.
- [2] B. Chen, M. Li, J. Wang, F.-X. Wu, Disease gene identification by using graph kernels and markov random fields, Science China. Life Sciences
205 57 (11) (2014) 1054.

	BioGPS							
Disease	DK	MD	MED	RL	CDNK1	CDNK2	CDNK3	CDNK4
1	51.9/8	57.4/7	59.0/6	59.2/5	65.1/4	69.5/2	69.3/3	70.3/1
2	81.7/5	78.5/6	75.2/7	75.0/8	88.3/2	88.8/1	85.1/4	86.8/3
3	64.3/7	59.6/8	71.6/3	71.8/2	65.5/5	72.5/1	64.7/6	66.4/4
4	65.3/7	58.2/8	67.8/6	67.8/5	71.9/4	78.7/1	73.9/3	77.0/2
5	64.0/8	64.1/7	66.5/5	66.2/6	75.9/4	76.2/3	76.9/2	77.4/1
6	74.6/5	70.2/8	71.0/7	71.2/6	79.3/2	83.7/1	76.7/4	79.0/3
7	73.0/5	66.7/8	75.4/3	75.6/2	68.8/6	73.9/4	67.3/7	76.9/1
8	74.4/8	76.8/3	76.2/5	76.4/4	74.7/7	77.7/1	76.0/6	76.8/2
9	71.5/2	65.6/8	67.7/5	69.9/3	66.8/7	71.7/1	67.1/6	68.1/4
10	54.0/6	50.3/8	56.1/5	51.1/7	77.6/4	82.7/1	80.5/2	80.0/3
11	58.2/7	51.3/8	59.3/6	59.3/5	71.8/4	80.2/1	75.3/3	77.1/2
\overline{AUC}	66.6	63.5	67.8	67.6	73.2	77.8	73.9	76.0
\overline{Rank}	6.18	7.18	5.27	4.82	4.45	1.55	4.18	2.36

Table 1: *Predictive performance on 11 gene-disease associations in percentage using network induced by the BioGPS. Best results are in bold. We report the AUC ROC and the rank for each kernel method in which CDNK1 uses ontology for discrete labels, CDNK2 uses ontology for both discrete and continuous labels, CDNK3 uses node degree for discrete labels and CDNK4 uses degree and ontology for both discrete and continuous labels, respectively.*

	Pathways							
Disease	DK	MD	MED	RL	CDNK1	CDNK2	CDNK3	CDNK4
1	74.7/8	76.4/7	78.7/6	78.8/5	80.2/3	82.4/1	80.6/2	79.4/4
2	55.1/8	64.9/7	76.6/6	76.6/5	81.1/2	80.3/3	79.8/4	82.2/1
3	55.0/8	62.7/7	64.1/5	65.6/4	67.1/3	63.6/6	68.5/2	71.5/1
4	54.3/8	65.2/7	73.7/1	73.7/2	66.1/5	68.1/3	67.5/4	65.6/6
5	52.9/8	55.7/7	62.7/5	62.7/6	68.3/2	69.6/1	66.3/4	66.8/3
6	83.4/8	92.7/7	96.5/2	96.5/1	93.0/6	94.1/5	94.4/4	95.5/3
7	84.5/8	88.3/7	89.4/3	89.5/2	88.5/6	88.5/5	88.7/4	90.5/1
8	53.7/8	65.6/7	72.0/6	72.3/5	72.5/3	72.5/2	72.3/4	76.5/1
9	52.5/8	64.9/5	64.2/7	64.2/6	81.3/1	81.0/2	79.6/3	78.8/4
10	68.8/6	65.4/8	74.4/5	74.4/4	66.9/7	76.8/2	76.1/3	79.5/1
11	53.7/8	69.2/7	74.6/5	74.1/6	77.0/2	78.7/1	75.4/4	77.0/3
\overline{AUC}	62.6	70.1	75.2	75.3	76.5	77.8	77.2	78.5
\overline{Rank}	7.82	6.91	4.64	4.18	3.64	2.82	3.45	2.55

Table 2: *Predictive performance on 11 gene-disease associations in percentage using network induced by the Pathways. Best results are in bold. We report the AUC ROC and the rank for each kernel method in which CDNK1 uses ontology for discrete labels, CDNK2 uses ontology for both discrete and continuous labels, CDNK3 uses node degree for discrete labels and CDNK4 uses degree and ontology for discrete and continuous labels, respectively.*

- [3] F. Fouss, L. Yen, A. Pirotte, M. Saerens, An experimental investigation of graph kernels on a collaborative recommendation task, in: Data Mining, 2006. ICDM'06. Sixth International Conference on, IEEE, 2006, pp. 863–868.
- [4] P. Chebotarev, E. Shamis, The matrix-forest theorem and measuring relations in small social groups, arXiv preprint math/0602070.
- [5] D. Haussler, Convolution kernels on discrete structures, Tech. rep., Technical report, Department of Computer Science, University of California at Santa Cruz (1999).
- [6] G. O. Consortium, et al., The gene ontology (go) database and informatics resource, Nucleic acids research 32 (suppl 1) (2004) D258–D261.
- [7] J. I. Alvarez-Hamelin, L. Dall'Asta, A. Barrat, A. Vespignani, k-core decomposition: A tool for the visualization of large scale networks, arXiv preprint cs/0504107.
- [8] R. E. Tarjan, Decomposition by clique separators, Discrete mathematics 55 (2) (1985) 221–232.
- [9] F. Costa, K. De Grave, Fast neighborhood subgraph pairwise distance kernel, in: Proceedings of the 26th International Conference on Machine Learning, Omnipress, 2010, pp. 255–262.
- [10] K.-I. Goh, M. E. Cusick, D. Valle, B. Childs, M. Vidal, A.-L. Barabási, The human disease network, Proceedings of the National Academy of Sciences 104 (21) (2007) 8685–8690.